

Project

December 5, 2025

1 MAIN OBJECTIVE

1.1 USING UNSUPERVISED AND SUPERVISED LEARNING MODELS TO SEGREGATE AND PREDICT THE CLASS OF PUMPKIN SEED

2 DATASET USED = PUMPKIN SEED DATASET

Pumpkin seeds are frequently consumed as confection worldwide because of their adequate amount of protein, fat, carbohydrate, and mineral contents. This study was carried out on the two most important and quality types of pumpkin seeds, ‘‘Urgup_Sivrisi’’ and ‘‘Cercevelik’’, generally grown in Urgup and Karacaoren regions in Turkey. However, morphological measurements of 2500 pumpkin seeds of both varieties were made possible by using the gray and binary forms of threshold techniques. Considering morphological features, all the data were modeled with five different machine learning methods: Logistic Regression (LR), Multilayer Perceptrons (MLP), Support Vector Machine (SVM) and Random Forest (RF), and k-Nearest Neighbor (k-NN), which further determined the most successful method for classifying pumpkin seed varieties. However, the performances of the models were determined with the help of the 10 kfold cross-validation method. The accuracy rates of the classifiers were obtained as LR 87.92 percent, MLP 88.52 percent, SVM 88.64 percent, RF 87.56 percent, and k-NN 87.64 percent. DATASET: https://www.muratkoklu.com/datasets/Pumpkin_Seeds_Dataset.zip

TOTAL ROWS = 2500

FEATURES: |FEATURE|N ROWS| TYPE| |——|——|——| |0 Area | 2500 non-null | int64| 1 Perimeter | 2500 non-null |float64| 2 Major_Axis_Length | 2500 non-null |float64| 3 Minor_Axis_Length | 2500 non-null | float64| 4 Convex_Area | 2500 non-null | int64 | 5 Equiv_Diameter | 2500 non-null |float64| 6 Eccentricity | 2500 non-null | float64| 7 Solidity | 2500 non-null | float64| 8 Extent | 2500 non-null | float64| 9 Roundness | 2500 non-null | float64| 10 Aspect_Ration | 2500 non-null | float64| 11 Compactness | 2500 non-null | float64| 12 Class | 2500 non-null | object | 13 LR | 2500 non-null | float64|

```
[40]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix,classification_report
from sklearn.cluster import KMeans, DBSCAN, HDBSCAN, MeanShift
```

```
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import GridSearchCV
```

```
[7]: dataset = pd.read_excel('Pumpkin_Seeds_Dataset.xlsx')
dataset.head()
```

```
[7]:    Area  Perimeter  Major_Axis_Length  Minor_Axis_Length  Convex_Area \
0   56276      888.242           326.1485            220.2388      56831
1   76631     1068.146           417.1932            234.2289      77280
2   71623     1082.987           435.8328            211.0457      72663
3   66458      992.051           381.5638            222.5322      67118
4   66107      998.146           383.8883            220.4545      67117

      Equiv_Diameter  Eccentricity  Solidity  Extent  Roundness  Aspect_Ration \
0          267.6805       0.7376    0.9902  0.7453     0.8963        1.4809
1          312.3614       0.8275    0.9916  0.7151     0.8440        1.7811
2          301.9822       0.8749    0.9857  0.7400     0.7674        2.0651
3          290.8899       0.8123    0.9902  0.7396     0.8486        1.7146
4          290.1207       0.8187    0.9850  0.6752     0.8338        1.7413

  Compactness      Class
0      0.8207  Çerçevevik
1      0.7487  Çerçevevik
2      0.6929  Çerçevevik
3      0.7624  Çerçevevik
4      0.7557  Çerçevevik
```

```
[8]: dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Area             2500 non-null   int64  
 1   Perimeter        2500 non-null   float64 
 2   Major_Axis_Length 2500 non-null   float64 
 3   Minor_Axis_Length 2500 non-null   float64 
 4   Convex_Area      2500 non-null   int64  
 5   Equiv_Diameter   2500 non-null   float64 
 6   Eccentricity     2500 non-null   float64 
 7   Solidity          2500 non-null   float64 
 8   Extent            2500 non-null   float64 
 9   Roundness          2500 non-null   float64 
 10  Aspect_Ration     2500 non-null   float64 
 11  Compactness        2500 non-null   float64 
 12  Class             2500 non-null   object 
```

```
dtypes: float64(10), int64(2), object(1)
memory usage: 254.0+ KB
```

```
[9]: dataset.isna().sum()
```

```
[9]: Area          0
Perimeter      0
Major_Axis_Length 0
Minor_Axis_Length 0
Convex_Area     0
Equiv_Diameter  0
Eccentricity    0
Solidity        0
Extent          0
Roundness       0
Aspect_Ration   0
Compactness     0
Class           0
dtype: int64
```

```
[10]: dataset.describe()
```

```
[10]:          Area  Perimeter  Major_Axis_Length  Minor_Axis_Length \
count  2500.000000  2500.000000  2500.000000  2500.000000
mean   80658.220800 1130.279015  456.601840  225.794921
std    13664.510228 109.256418   56.235704  23.297245
min    47939.000000  868.485000  320.844600  152.171800
25%    70765.000000 1048.829750  414.957850  211.245925
50%    79076.000000 1123.672000  449.496600  224.703100
75%    89757.500000 1203.340500  492.737650  240.672875
max   136574.000000 1559.450000  661.911300  305.818000

          Convex_Area  Equiv_Diameter  Eccentricity  Solidity  Extent \
count  2500.000000  2500.000000  2500.000000  2500.000000  2500.000000
mean   81508.084400  319.334230  0.860879  0.989492  0.693205
std    13764.092788  26.891920  0.045167  0.003494  0.060914
min    48366.000000  247.058400  0.492100  0.918600  0.468000
25%    71512.000000  300.167975  0.831700  0.988300  0.658900
50%    79872.000000  317.305350  0.863700  0.990300  0.713050
75%    90797.750000  338.057375  0.897025  0.991500  0.740225
max   138384.000000  417.002900  0.948100  0.994400  0.829600

          Roundness  Aspect_Ration  Compactness
count  2500.000000  2500.000000  2500.000000
mean   0.791533    2.041702    0.704121
std    0.055924    0.315997    0.053067
min    0.554600    1.148700    0.560800
25%    0.751900    1.801050    0.663475
```

50%	0.797750	1.984200	0.707700
75%	0.834325	2.262075	0.743500
max	0.939600	3.144400	0.904900

3 FEATURE ENGINEERING

TAKE LATUS RECTUM EQUAL TO $2 * (\text{MINOR_AXIS_LENGTH})^2 / \text{MAJOR_AXIS_LENGTH}$
 # DATA CLEANING REMOVING SOLIDITY AS VARIANCE IN SOLIDITY VERY LITTLE
 NOT PROVIDING VALUABLE INFO

```
[11]: dataset['LR'] = 2*dataset['Major_Axis_Length']**2 / dataset['Minor_Axis_Length']
```

```
[12]: dataset.head()
```

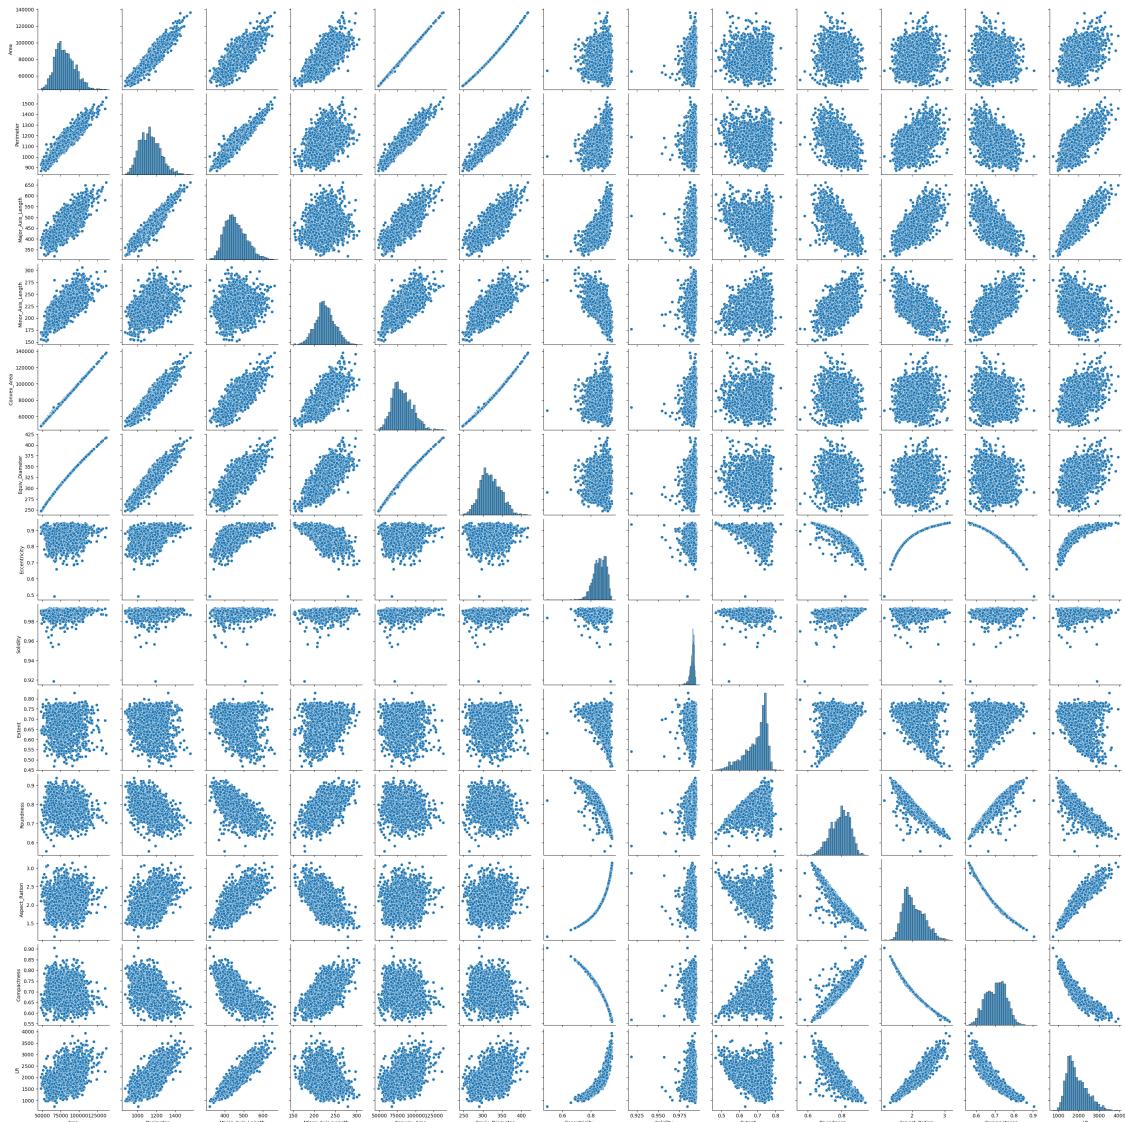
	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area	\\
0	56276	888.242	326.1485	220.2388	56831	
1	76631	1068.146	417.1932	234.2289	77280	
2	71623	1082.987	435.8328	211.0457	72663	
3	66458	992.051	381.5638	222.5322	67118	
4	66107	998.146	383.8883	220.4545	67117	

	Equiv_Diameter	Eccentricity	Solidity	Extent	Roundness	Aspect_Ration	\\
0	267.6805	0.7376	0.9902	0.7453	0.8963	1.4809	
1	312.3614	0.8275	0.9916	0.7151	0.8440	1.7811	
2	301.9822	0.8749	0.9857	0.7400	0.7674	2.0651	
3	290.8899	0.8123	0.9902	0.7396	0.8486	1.7146	
4	290.1207	0.8187	0.9850	0.6752	0.8338	1.7413	

	Compactness	Class	LR
0	0.8207	Çerçeveilik	965.977331
1	0.7487	Çerçeveilik	1486.154494
2	0.6929	Çerçeveilik	1800.086233
3	0.7624	Çerçeveilik	1308.493184
4	0.7557	Çerçeveilik	1336.967282

```
[32]: import seaborn as sns
sns.pairplot(dataset)
```

```
[32]: <seaborn.axisgrid.PairGrid at 0x70e765b76f90>
```



```
[13]: x = dataset.drop(['Major_Axis_Length', 'Class', "Solidity", "Minor_Axis_Length"], axis=1)
y = dataset['Class']
```

```
[24]: np.unique(y)
```

```
[24]: array([0, 1])
```

ENCODING CLASS

```
[15]: from sklearn.preprocessing import LabelEncoder
```

```
[17]: le = LabelEncoder()
y = le.fit_transform(y)

[18]: y

[18]: array([0, 0, 0, ..., 1, 1, 1], shape=(2500,))

[28]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2)

[20]: svc = SVC()
```

4 SVM

Achieved Accuracy of 87.6% on test dataset

```
[86]: from sklearn.pipeline import Pipeline
pipe = Pipeline([
    ('scaler', StandardScaler()),
    ('classifier', SVC())
])

[87]: pipe.fit(x_train,y_train)

[87]: Pipeline(steps=[('scaler', StandardScaler()), ('classifier', SVC())])

[88]: accuracy_score(pipe.predict(x_test), y_test)

[88]: 0.876
```

5 Random Forest Classifier

Accuracy: 87.6

```
[35]: from sklearn.ensemble import RandomForestClassifier

[71]: pipe2= Pipeline([
    ('scaler', StandardScaler()),
    ('tree', RandomForestClassifier())
])

[72]: pipe2.fit(x_train, y_train)
accuracy_score(pipe2.predict(x_test), y_test)*100

[72]: 87.6
```

6 Decision Tree

Accuracy 84.2

```
[79]: from sklearn.tree import DecisionTreeClassifier
pipe3 = Pipeline([
    ('scaler', StandardScaler()),
    ('tree', DecisionTreeClassifier())

])
pipe3.fit(x_train, y_train)
accuracy_score(pipe3.predict(x_test), y_test)*100
```

[79]: 84.2

```
[74]: from sklearn.mixture import GaussianMixture
from sklearn.decomposition import PCA, KernelPCA
```

6.0.1 USING DIMENSIONALITY REDUCTION

PCA WITH RANDOM TREE CLASSIFIER Accuracy 85.4

```
[99]: pipe4 = Pipeline([
    ('scaler', StandardScaler()),
    ('pca', PCA(n_components=2)),
    ('tree', RandomForestClassifier())

])
pipe4.fit(x_train, y_train)
accuracy_score(pipe4.predict(x_test), y_test)*100
```

[99]: 85.39999999999999

7 Model Evaluation

Random Forest and SVM are better than Decision Tree classifier PCA is not helping improving the model performance. For improving accuracy we can shift to Deep Neural Networks

8

8.1 UNSUPERVISED IS NOT RECOMMENDED ON THIS DATASET

ACCURACY IS AROUND 50 - 55%

```
[132]: km = KMeans(n_clusters=2)
```

```
[133]: km.fit(x_train_transformed, y_train)
```

```
[133]: KMeans(n_clusters=2)
```

```
[134]: pca = PCA(n_components=3)
pca.fit(x_train)
```

```
[134]: PCA(n_components=3)
```

```
[135]: preds = km.predict(pca.transform(x_test))
```

```
[136]: accuracy_score(preds, y_test)
```

```
[136]: 0.566
```

8.1.1 INSIGHTS

MODEL CAN BE DEPLOYED AT LARGE SCALE FARM IN SEGREGATING PUMPKIN SEEDS